

```
!pip install -q hvplot
# Install memory_profiler
%pip install memory_profiler

# Load memory_profiler extension
%load_ext memory_profiler

^C
Requirement already satisfied: memory_profiler in c:\users\fahad\anaconda3\envs\sea600\lib\site-packages (0.58.0)
Requirement already satisfied: psutil in c:\users\fahad\anaconda3\envs\sea600\lib\site-packages (from memory_profiler) (5.9.0)
Note: you may need to restart the kernel to use updated packages.
The memory_profiler extension is already loaded. To reload it, use:
    %reload_ext memory_profiler
```

```
import pandas as pd
import numpy as np
import seaborn as sns
from scipy import stats
import matplotlib.pyplot as plt
import hvplot.pandas
from IPython.display import display

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix
from sklearn.model_selection import cross_val_score
from sklearn.metrics import make_scorer, f1_score, roc_auc_score
from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import roc_curve, auc, precision_recall_curve, average_precision_score, RocCurveDisplay
from sklearn.model_selection import StratifiedKFold
from sklearn.base import clone
from sklearn.discriminant_analysis import QuadraticDiscriminantAnalysis
from sklearn.svm import SVC
from sklearn.discriminant_analysis import LinearDiscriminantAnalysis
from sklearn.tree import DecisionTreeClassifier

import time

pd.set_option('display.float', '{:.2f}'.format)
pd.set_option('display.max_columns', 50)
pd.set_option('display.max_rows', 50)
```

```
""  
""  
""
```

Introduction

- LendingClub is a US peer-to-peer lending company, headquartered in San Francisco, California.
- What we are trying to do is to predict the loan approval of a person based on the data we have been given
- Lets now **load** the data and **describe** it.

```
data = pd.read_csv("train_lending_club.csv")  
print("Data imported successfully")
```

```
data.head()
```

Data imported successfully

	issue_d	sub_grade	term	home_ownership	fico_range_low
total_acc \					
0	2014-01-01	D1	60 months	MORTGAGE	660.00
18.00					
1	2014-01-01	C4	60 months	MORTGAGE	740.00
26.00					
2	2014-01-01	A4	36 months	RENT	700.00
47.00					
3	2014-01-01	D1	60 months	MORTGAGE	665.00
26.00					
4	2014-01-01	C3	36 months	MORTGAGE	700.00
15.00					

	pub_rec	revol_util	annual_inc	int_rate	dti	purpose
\						
0	0.00	86.80	40440.00	16.99	15.16	credit_card
1	0.00	103.50	59000.00	15.61	16.74	credit_card
2	0.00	11.40	40000.00	7.90	20.34	debt_consolidation
3	0.00	56.20	70000.00	16.99	23.15	debt_consolidation
4	1.00	67.10	120000.00	14.98	17.88	vacation

	mort_acc	loan_amnt	application_type	installment
verification_status \				
0	1.00	17775.00	Individual	441.66
Verified				

1	4.00	29175.00	Individual	703.45	
Verified					
2	1.00	6000.00	Individual	187.75	Not
Verified					
3	7.00	15600.00	Individual	387.62	
Verified					
4	1.00	10000.00	Individual	346.56	
Verified					

	pub_rec_bankruptcies	addr_state	initial_list_status
fico_range_high \			
0	0.00	AR	f
664.00			
1	0.00	VT	f
744.00			
2	0.00	TX	w
704.00			
3	0.00	IN	f
669.00			
4	0.00	LA	f
704.00			

	revol_bal	id	open_acc	emp_length	loan_status \
0	17264.00	11024793	11.00	2.00	1
1	6725.00	10596078	8.00	3.00	1
2	7613.00	10775616	27.00	1.00	1
3	14173.00	10765610	9.00	11.00	1
4	2549.00	10794837	8.00	2.00	1

	time_to_earliest_cr_line
0	478656.00
1	541728.00
2	657590.40
3	328838.40
4	305164.80

Description of the data

```
data.describe()
```

	fico_range_low	total_acc	pub_rec	revol_util	annual_inc
int_rate \					
count	236846.00	236846.00	236846.00	236846.00	236846.00
mean	694.35	25.57	0.24	52.85	77379.53
std	30.93	12.20	0.67	24.01	77290.67
min	660.00	2.00	0.00	0.00	0.00

25%	670.00	17.00	0.00	35.00	46000.00
9.17					
50%	685.00	24.00	0.00	53.00	65000.00
12.29					
75%	710.00	32.00	0.00	71.20	92000.00
15.31					
max	845.00	169.00	86.00	182.80	9573072.00
30.99					

	dti	mort_acc	loan_amnt	installment
pub_rec_bankruptcies \				
count	236846.00	236846.00	236846.00	236846.00
236846.00				
mean	18.53	1.73	14626.43	441.92
0.14				
std	9.21	2.04	8611.59	257.21
0.39				
min	0.00	0.00	1000.00	14.77
0.00				
25%	12.05	0.00	8000.00	254.07
0.00				
50%	17.89	1.00	12450.00	378.20
0.00				
75%	24.52	3.00	20000.00	588.37
0.00				
max	999.00	47.00	40000.00	1534.88
9.00				

	fico_range_high	revol_bal	id	open_acc	emp_length \
count	236846.00	236846.00	236846.00	236846.00	236846.00
mean	698.35	16956.26	49269391.36	11.81	5.99
std	30.93	23829.01	20980382.07	5.59	4.24
min	664.00	0.00	56705.00	0.00	0.00
25%	674.00	6127.00	33191497.25	8.00	2.00
50%	689.00	11399.00	52979215.50	11.00	6.00
75%	714.00	20431.00	66645529.25	15.00	11.00
max	850.00	2560703.00	84363456.00	76.00	11.00

	loan_status	time_to_earliest_cr_line
count	236846.00	236846.00
mean	0.83	526420.51
std	0.37	241312.11
min	0.00	97113.60
25%	1.00	360374.40
50%	1.00	478483.20
75%	1.00	647049.60
max	1.00	2240524.80

data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 236846 entries, 0 to 236845
Data columns (total 27 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   issue_d                               236846 non-null  object
1   sub_grade                             236846 non-null  object
2   term                                  236846 non-null  object
3   home_ownership                       236846 non-null  object
4   fico_range_low                       236846 non-null  float64
5   total_acc                            236846 non-null  float64
6   pub_rec                              236846 non-null  float64
7   revol_util                           236846 non-null  float64
8   annual_inc                           236846 non-null  float64
9   int_rate                             236846 non-null  float64
10  dti                                   236846 non-null  float64
11  purpose                              236846 non-null  object
12  mort_acc                             236846 non-null  float64
13  loan_amnt                            236846 non-null  float64
14  application_type                     236846 non-null  object
15  installment                          236846 non-null  float64
16  verification_status                 236846 non-null  object
17  pub_rec_bankruptcies                236846 non-null  float64
18  addr_state                           236846 non-null  object
19  initial_list_status                 236846 non-null  object
20  fico_range_high                     236846 non-null  float64
21  revol_bal                            236846 non-null  float64
22  id                                   236846 non-null  int64
23  open_acc                             236846 non-null  float64
24  emp_length                           236846 non-null  float64
25  loan_status                          236846 non-null  int64
26  time_to_earliest_cr_line             236846 non-null  float64
dtypes: float64(16), int64(2), object(9)
memory usage: 48.8+ MB

```

Data Preprocessing

- Remove or fill any missing data
- scaling
- convert categorical data into numeric

```
print(f"The Length of the data: {data.shape}")
```

```
The Length of the data: (236846, 27)
```

```
# Missing values
```

```

for column in data.columns:
    if data[column].isna().sum() != 0:
        missing = data[column].isna().sum()
        portion = (missing / data.shape[0]) * 100

```

```

        print(f"{column}": number of missing values '{missing}' ==>
'{portion:.3f}%')
data.emp_length.unique()
array([ 2.,  3.,  1., 11.,  9.,  8.,  6.,  4.,  0.,  5.,  7.])

for year in data.emp_length.unique():
    print(f"{year} years in this position:")
    print(f"{data[data.emp_length ==
year].loan_status.value_counts(normalize=True)}")
    print('=====')

2.0 years in this position:
loan_status
1    0.84
0    0.16
Name: proportion, dtype: float64
=====
3.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
1.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
11.0 years in this position:
loan_status
1    0.84
0    0.16
Name: proportion, dtype: float64
=====
9.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
8.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
6.0 years in this position:

```

```

loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
4.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
0.0 years in this position:
loan_status
1    0.80
0    0.20
Name: proportion, dtype: float64
=====
5.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====
7.0 years in this position:
loan_status
1    0.83
0    0.17
Name: proportion, dtype: float64
=====

```

Label Encoding

- Turning Categorical values into numerical

```

# List of categorical columns to convert
categorical_columns = ['sub_grade', 'term', 'home_ownership',
                       'purpose', 'application_type', 'verification_status',
                       'initial_list_status']

# One-hot encode these columns
data = pd.get_dummies(data, columns=categorical_columns)

# remove issue date because before hand we don't know when the loan
will be issued
# drop loan_status because it is the target
# drop id, addr_state because it is not useful

# Store 'loan_status' in a separate variable Y
y_train = data['loan_status']

```

```
# Now we can safely drop 'issue_d', 'sub_grade', 'loan_status', and
'id' from the dataset
data = data.drop(['issue_d', 'loan_status', 'id', 'addr_state'],
axis=1)
X_train = data
```

```
# 'X_train' now contains only the features, and 'y_train' is your
target variable
```

```
X_train.head()
```

	fico_range_low	total_acc	pub_rec	revol_util	annual_inc
0	660.00	18.00	0.00	86.80	40440.00
1	740.00	26.00	0.00	103.50	59000.00
2	700.00	47.00	0.00	11.40	40000.00
3	665.00	26.00	0.00	56.20	70000.00
4	700.00	15.00	1.00	67.10	120000.00

	mort_acc	loan_amnt	installment	pub_rec_bankruptcies
0	1.00	17775.00	441.66	0.00
1	4.00	29175.00	703.45	0.00
2	1.00	6000.00	187.75	0.00
3	7.00	15600.00	387.62	0.00
4	1.00	10000.00	346.56	0.00

	revol_bal	open_acc	emp_length	time_to_earliest_cr_line
0	17264.00	11.00	2.00	478656.00
1	6725.00	8.00	3.00	541728.00
2	7613.00	27.00	1.00	657590.40
3	14173.00	9.00	11.00	328838.40
4	2549.00	8.00	2.00	305164.80

	sub_grade_A2	sub_grade_A3	sub_grade_A4	sub_grade_A5
sub_grade_B1 \				
0	False	False	False	False
False				
1	False	False	False	False
False				
2	False	False	True	False
False				
3	False	False	False	False
False				
4	False	False	False	False
False				
	sub_grade_B2	sub_grade_B3	sub_grade_B4	... term_ 60 months \
0	False	False	False	... True
1	False	False	False	... True
2	False	False	False	... False
3	False	False	False	... True
4	False	False	False	... False
	home_ownership_MORTGAGE	home_ownership_OWN	home_ownership_RENT	\
0		True	False	False
1		True	False	False
2		False	False	True
3		True	False	False
4		True	False	False
	purpose_car	purpose_credit_card	purpose_debt_consolidation	\
0	False	True		False
1	False	True		False
2	False	False		True
3	False	False		True
4	False	False		False
	purpose_educational	purpose_home_improvement	purpose_house	\
0	False	False	False	False
1	False	False	False	False
2	False	False	False	False
3	False	False	False	False
4	False	False	False	False
	purpose_major_purchase	purpose_medical	purpose_moving	
purpose_other \				
0	False	False	False	
False				
1	False	False	False	
False				
2	False	False	False	
False				
3	False	False	False	

False			
4	False	False	False
False			
	purpose_renewable_energy	purpose_small_business	purpose_vacation
\			
0	False	False	False
1	False	False	False
2	False	False	False
3	False	False	False
4	False	False	True
	purpose_wedding	application_type_Individual	
application_type_Joint App	\		
0	False	True	
False			
1	False	True	
False			
2	False	True	
False			
3	False	True	
False			
4	False	True	
False			
	verification_status_Not Verified	verification_status_Source	
Verified	\		
0	False		
False			
1	False		
False			
2	True		
False			
3	False		
False			
4	False		
False			
	verification_status_Verified	initial_list_status_f	
initial_list_status_w			
0	True	True	
False			
1	True	True	
False			
2	False	False	

True		
3	True	True
False		
4	True	True
False		

[5 rows x 77 columns]

Logistic Regression

```
# Scale the numeric features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)

logistic_classifier = LogisticRegression(max_iter=1000) # You can
adjust regularization parameters if needed
logistic_classifier.fit(X_train_scaled, y_train)

LogisticRegression(max_iter=1000)
```

KNN CLASSIFIER

```
# Train KNN Classifier
# in terminal set LOKY_MAX_CPU_COUNT= <number of cores you want to
use>

knn_classifier = KNeighborsClassifier(n_neighbors=5) # You can adjust
the number of neighbors (k) as needed
knn_classifier.fit(X_train_scaled, y_train)

KNeighborsClassifier()
```

Performance evaluation Milestone 1

```
# F1 Score for Logistic Regression
logistic_f1_scores = cross_val_score(logistic_classifier,
X_train_scaled, y_train, cv=5, scoring='f1')
print("Logistic Regression CV F1 Score:", logistic_f1_scores.mean())

# AUC for Logistic Regression
logistic_auc_scores = cross_val_score(logistic_classifier,
X_train_scaled, y_train, cv=5, scoring='roc_auc')
print("Logistic Regression CV AUC:", logistic_auc_scores.mean())

# Repeat for KNN
knn_f1_scores = cross_val_score(knn_classifier, X_train_scaled,
y_train, cv=5, scoring='f1')
print("KNN CV F1 Score:", knn_f1_scores.mean())

knn_auc_scores = cross_val_score(knn_classifier, X_train_scaled,
```

```
y_train, cv=5, scoring='roc_auc')
print("KNN CV AUC:", knn_auc_scores.mean())
```

Logistic Regression CV F1 Score: 0.9033793225069063
Logistic Regression CV AUC: 0.7789966094200176
KNN CV F1 Score: 0.8945972841418024
KNN CV AUC: 0.6943401702263249

ROC Curve for Logistic Regression

```
# Initialize StratifiedKFold
cv = StratifiedKFold(n_splits=5)

tprs = []
aucs = []
mean_fpr = np.linspace(0, 1, 100)

fig, ax = plt.subplots()
for i, (train, test) in enumerate(cv.split(X_train_scaled, y_train)):
    # Fit the logistic regression model on the training data
    logistic_classifier.fit(X_train_scaled[train], y_train[train])

    # Predict probabilities for the test set
    probas_ = logistic_classifier.predict_proba(X_train_scaled[test])

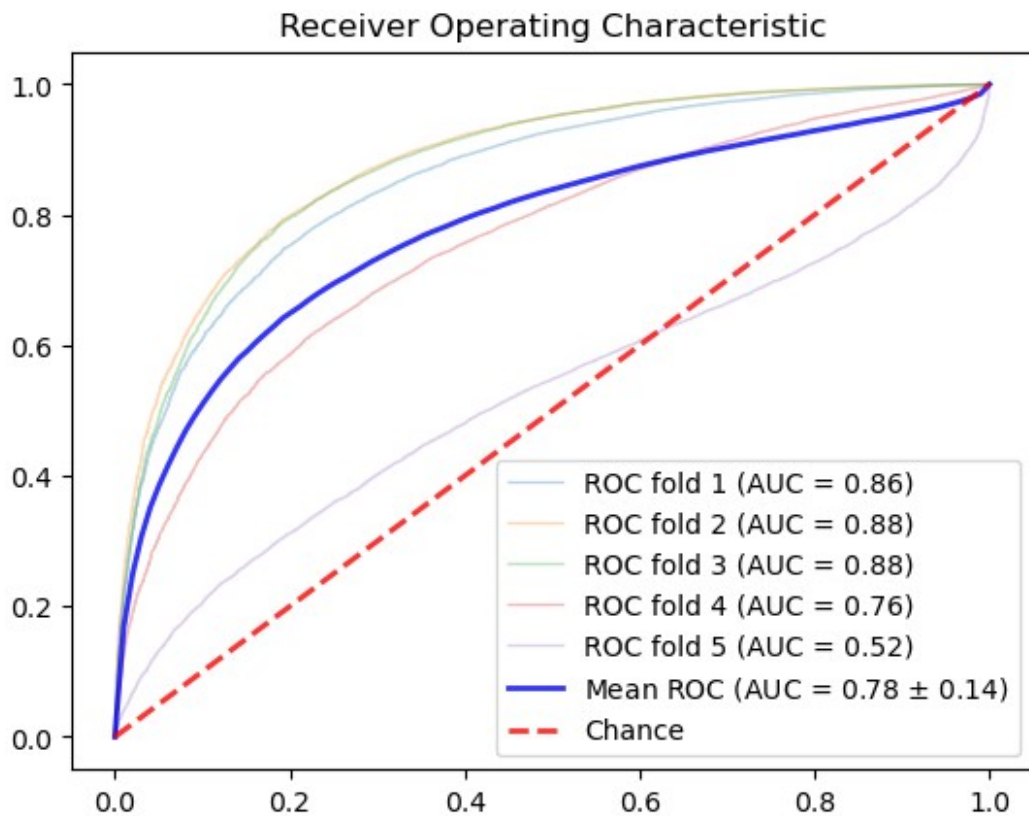
    # Compute ROC curve and area the curve
    fpr, tpr, thresholds = roc_curve(y_train[test], probas_[ :, 1])
    tprs.append(np.interp(mean_fpr, fpr, tpr))
    tprs[-1][0] = 0.0
    roc_auc = auc(fpr, tpr)
    aucs.append(roc_auc)
    ax.plot(fpr, tpr, lw=1, alpha=0.3, label=f'ROC fold {i+1} (AUC = {roc_auc:.2f})')

# Plot the mean ROC curve
mean_tpr = np.mean(tprs, axis=0)
mean_tpr[-1] = 1.0
mean_auc = auc(mean_fpr, mean_tpr)
std_auc = np.std(aucs)
ax.plot(mean_fpr, mean_tpr, color='b', label=r'Mean ROC (AUC = %0.2f $\pm$ %0.2f)' % (mean_auc, std_auc),
        lw=2, alpha=.8)

# Plot the chance line
ax.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r',
        label='Chance', alpha=.8)

# Final plot adjustments
ax.set(xlim=[-0.05, 1.05], ylim=[-0.05, 1.05], title="Receiver Operating Characteristic")
```

```
ax.legend(loc="lower right")
plt.show()
```



```
def plot_cv_roc_curve(classifier, X, y, cv_splits=5,
                      classifier_name='Classifier'):
    """
    Plots ROC curves for each fold in cross-validation and the mean
    ROC curve.

    Parameters:
    - classifier: The classifier to use.
    - X: Feature set.
    - y: Target variable.
    - cv_splits: Number of cross-validation splits.
    - classifier_name: Name of the classifier (for legend labeling).
    """
    cv = StratifiedKFold(n_splits=cv_splits)
    tprs = []
    aucs = []
    mean_fpr = np.linspace(0, 1, 100)
    fig, ax = plt.subplots()

    for i, (train, test) in enumerate(cv.split(X, y)):
```

```

        classifier.fit(X[train], y[train])
        probas_ = classifier.predict_proba(X[test])
        fpr, tpr, thresholds = roc_curve(y[test], probas_[ :, 1])
        tprs.append(np.interp(mean_fpr, fpr, tpr))
        tprs[-1][0] = 0.0
        roc_auc = auc(fpr, tpr)
        aucs.append(roc_auc)
        ax.plot(fpr, tpr, lw=1, alpha=0.3, label=f'Fold {i+1}
{classifier_name} (AUC = {roc_auc:.2f})')

    mean_tpr = np.mean(tprs, axis=0)
    mean_tpr[-1] = 1.0
    mean_auc = auc(mean_fpr, mean_tpr)
    std_auc = np.std(aucs)
    ax.plot(mean_fpr, mean_tpr, color='b', label=f'Mean
{classifier_name} (AUC = {mean_auc:.2f} $\pm$ {std_auc:.2f})', lw=2,
alpha=.8)
    ax.plot([0, 1], [0, 1], linestyle='--', lw=2, color='r',
label='Chance', alpha=.8)
    ax.set(xlim=[-0.05, 1.05], ylim=[-0.05, 1.05], title=f"Receiver
Operating Characteristic - {classifier_name}")
    ax.legend(loc="lower right")
    plt.show()

# Measure CPU time for training KNN Classifier
start_time = time.time()
knn_classifier.fit(X_train_scaled, y_train)
end_time = time.time()
print("Time taken for training KNN Classifier:", end_time -
start_time, "seconds")

# Measure memory usage for training KNN Classifier
%memit knn_classifier.fit(X_train_scaled, y_train)

# Measure CPU time for training Logistic Regression Classifier
start_time = time.time()
logistic_classifier.fit(X_train_scaled, y_train)
end_time = time.time()
print("Time taken for training Logistic Classifier:", end_time -
start_time, "seconds")

# Measure memory usage for training KNN Classifier
%memit logistic_classifier.fit(X_train_scaled, y_train)

Time taken for training KNN Classifier: 0.08687520027160645 seconds
peak memory: 763.15 MiB, increment: 142.79 MiB
Time taken for training Logistic Classifier: 1.1194474697113037
seconds
peak memory: 768.02 MiB, increment: 147.64 MiB

```

Linear Discriminant Analysis AND Decision Trees

```
# Define the LDA model
lda_model = LinearDiscriminantAnalysis()
# Measure training time for LDA
start_time = time.time()
lda_model.fit(X_train_scaled, y_train)
lda_training_time = time.time() - start_time

print(f"LDA Training Time: {lda_training_time} seconds")

# Define the Decision Tree model
dt_model = DecisionTreeClassifier()

# Measure training time for Decision Trees
start_time = time.time()
dt_model.fit(X_train_scaled, y_train)
dt_training_time = time.time() - start_time

print(f"Decision Tree Training Time: {dt_training_time} seconds")
```

LDA Training Time: 2.908019542694092 seconds
Decision Tree Training Time: 15.742506742477417 seconds

Performance and Resource Utilization

```
# AUC and F1 for LDA
lda_auc = cross_val_score(lda_model, X_train_scaled, y_train, cv=5,
scoring='roc_auc').mean()
lda_f1 = cross_val_score(lda_model, X_train_scaled, y_train, cv=5,
scoring='f1').mean()

print(f"LDA Mean AUC: {lda_auc}")
print(f"LDA Mean F1: {lda_f1}")

# AUC and F1 for Decision Trees
dt_auc = cross_val_score(dt_model, X_train_scaled, y_train, cv=5,
scoring='roc_auc').mean()
dt_f1 = cross_val_score(dt_model, X_train_scaled, y_train, cv=5,
scoring='f1').mean()

print(f"Decision Tree Mean AUC: {dt_auc}")
print(f"Decision Tree Mean F1: {dt_f1}")
```


LDA Mean AUC: 0.7762112146124798
LDA Mean F1: 0.9008621176642944
Decision Tree Mean AUC: 0.5538783998959265
Decision Tree Mean F1: 0.8094133830039187

```
# For Logistic Regression
plot_cv_roc_curve(logistic_classifier, X_train_scaled, y_train,
classifier_name='Logistic Regression')

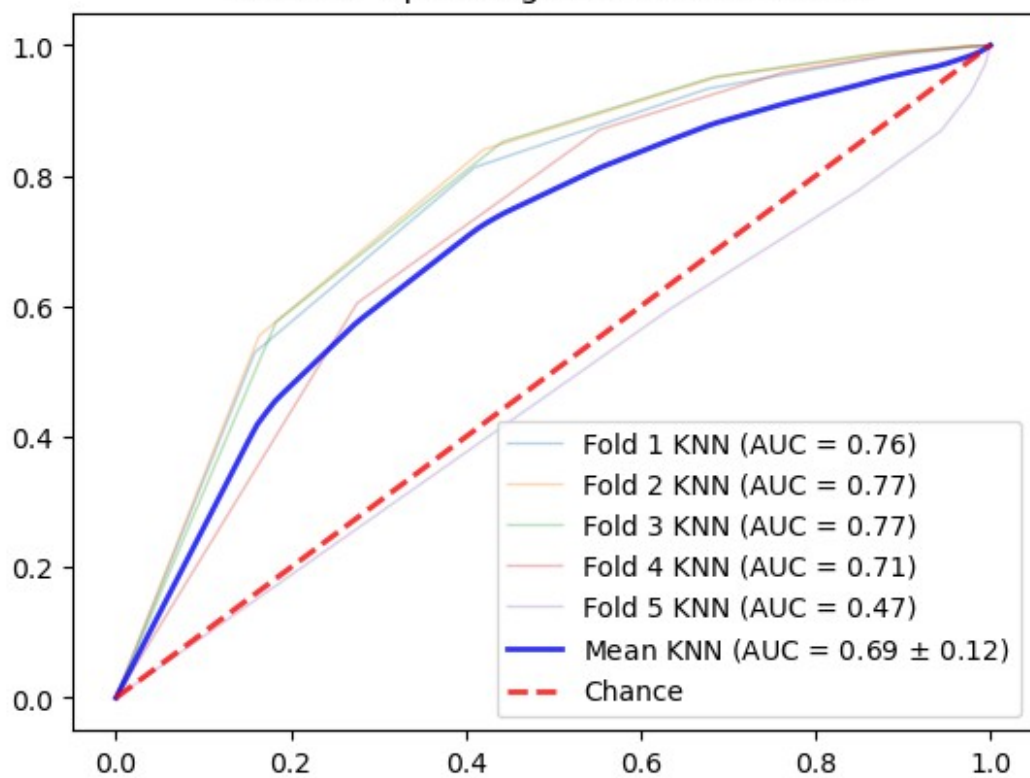
# For KNN
plot_cv_roc_curve(knn_classifier, X_train_scaled, y_train,
classifier_name='KNN')

# For LDA
plot_cv_roc_curve(lda_model, X_train_scaled, y_train,
classifier_name='LDA')

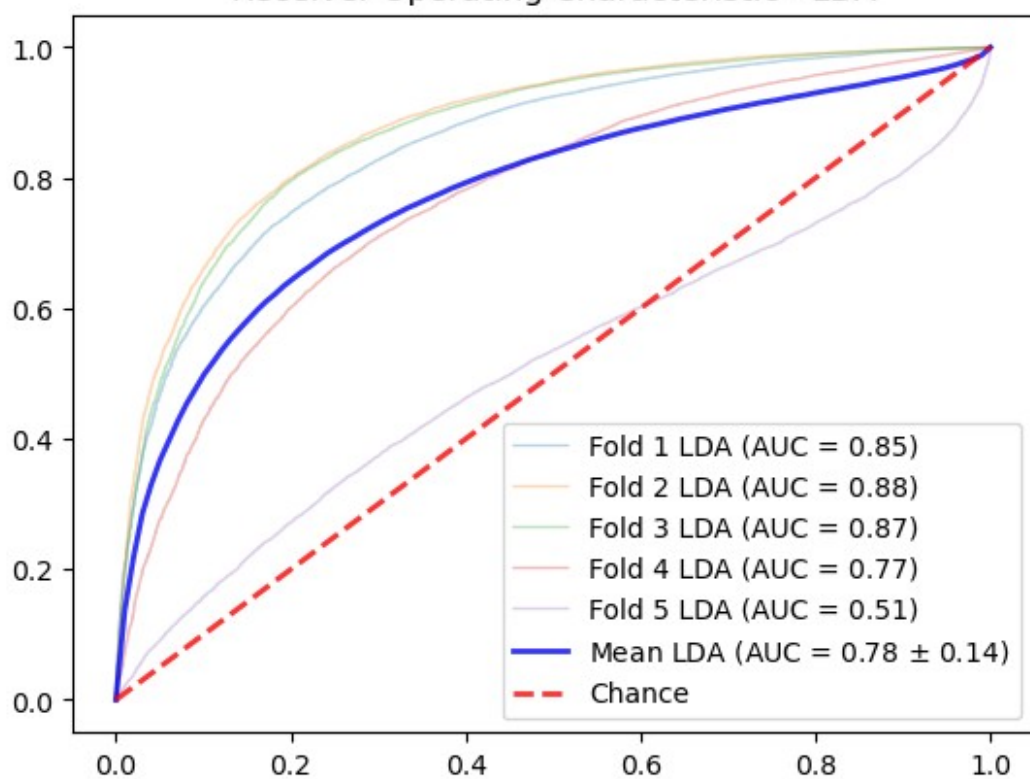
# For Decision Trees
plot_cv_roc_curve(dt_model, X_train_scaled, y_train,
classifier_name='Decision Tree')
```




Receiver Operating Characteristic - KNN



Receiver Operating Characteristic - LDA





Precision Recall Curve

```
def plot_cv_precision_recall_curve(classifier, X, y, cv_splits=5,
    classifier_name='Classifier'):
    """
    Plots Precision-Recall curves for each fold in cross-validation
    and the mean Precision-Recall curve.

    Parameters:
    - classifier: The classifier to use.
    - X: Feature set.
    - y: Target variable.
    - cv_splits: Number of cross-validation splits.
    - classifier_name: Name of the classifier (for legend labeling).
    """
    cv = StratifiedKFold(n_splits=cv_splits)
    mean_recall = np.linspace(0, 1, 100)
    precisions = []
    aucs = []
    fig, ax = plt.subplots()

    for i, (train, test) in enumerate(cv.split(X, y)):
        classifier.fit(X[train], y[train])
        probas_ = classifier.predict_proba(X[test])
        precision, recall, thresholds =
```

```

precision_recall_curve(y[test], probas[:, 1])
    rev_recall = recall[::-1] # Reverse the recall array
    rev_precision = precision[::-1] # Reverse the precision array
    precisions.append(np.interp(mean_recall, rev_recall,
rev_precision))
    precisions[-1][0] = 1.0
    pr_auc = auc(rev_recall, rev_precision)
    aucs.append(pr_auc)
    ax.plot(recall, precision, lw=1, alpha=0.3, label=f'Fold {i+1}
{classifier_name} (AUC = {pr_auc:.2f})')

    mean_precision = np.mean(precisions, axis=0)
    mean_auc = auc(mean_recall, mean_precision)
    std_auc = np.std(aucs)
    ax.plot(mean_recall, mean_precision, color='b', label=f'Mean
{classifier_name} (AUC = {mean_auc:.2f}  $\pm$  {std_auc:.2f})', lw=2,
alpha=.8)
    ax.set(xlim=[0, 1.05], ylim=[0, 1.05], title=f"Precision-Recall
Curve - {classifier_name}")
    ax.legend(loc="lower left")
    plt.xlabel('Recall')
    plt.ylabel('Precision')
    plt.show()

# For Logistic Regression
plot_cv_precision_recall_curve(logistic_classifier, X_train_scaled,
y_train, classifier_name='Logistic Regression')

# For KNN
plot_cv_precision_recall_curve(knn_classifier, X_train_scaled,
y_train, classifier_name='KNN')

# For LDA
plot_cv_precision_recall_curve(lda_model, X_train_scaled, y_train,
classifier_name='LDA')

# For Decision Trees
plot_cv_precision_recall_curve(dt_model, X_train_scaled, y_train,
classifier_name='Decision Tree')

```

